

# Heuristic (Informed) Search

(Where we try to choose smartly)

R&N: Chap. 4, Sect. 4.1-3

Recall that the ordering  
of FRINGE defines the  
search strategy

## Search Algorithm #2

SEARCH#2

1. INSERT(initial-node, FRINGE)
2. Repeat:
  - a. If empty(FRINGE) then return failure
  - b.  $N \leftarrow \text{REMOVE}(\text{FRINGE})$
  - c.  $s \leftarrow \text{STATE}(N)$
  - d. If GOAL?( $s$ ) then return path or goal state
  - e. For every state  $s'$  in SUCCESSORS( $s$ )
    - i. Create a node  $N'$  as a successor of  $N$
    - ii. INSERT( $N'$ , FRINGE)

# Best-First Search

- It exploits **state description** to estimate how “good” each search node is
- An **evaluation function**  $f$  maps each node  $N$  of the search tree to a real number  $f(N) \geq 0$   
[Traditionally,  $f(N)$  is an estimated cost; so, the smaller  $f(N)$ , the more promising  $N$ ]
- **Best-first search** sorts the FRINGE in increasing  $f$   
[Arbitrary order is assumed among nodes with equal  $f$ ]

# Best-First Search

- It exploits state description to estimate how "good" each search node is

- An evaluation function  $f$  maps each node  $N$  of the search tree to a real number  $f(N) \geq 0$

[Traditionally,  $f(N)$  is the cost of the path from the root to  $N$ , the more precise the better]

"Best" does not refer to the quality of the generated path  
Best-first search does not generate optimal paths in general

- Best-first search sorts the FRINGE in increasing  $f$   
[Arbitrary order is assumed among nodes with equal  $f$ ]

# How to construct $f$ ?

- Typically,  $f(N)$  estimates:

- either the **cost of a solution path through  $N$**

Then  $f(N) = g(N) + h(N)$ , where

- $g(N)$  is the cost of the path from the initial node to  $N$
- $h(N)$  is an estimate of the cost of a path from  $N$  to a goal node

- or the **cost of a path from  $N$  to a goal node**

Then  $f(N) = h(N) \rightarrow$  **Greedy best-search**

- But there are no limitations on  $f$ . Any function of your choice is acceptable.  
But will it help the search algorithm?

# How to construct $f$ ?

- Typically,  $f(N)$  estimates:

- either the cost of a solution path through  $N$

Then  $f(N) = g(N) + h(N)$ , where

- $g(N)$  is the cost of the path from the initial node to  $N$
- $h(N)$  is an estimate of the cost of a path from  $N$  to a goal node

- or the cost of a path from  $N$  to a goal node

Then  $f(N) = h(N)$

Heuristic function



- But there are no limitations on  $f$ . Any function of your choice is acceptable.  
But will it help the search algorithm?

# Heuristic Function

- The **heuristic function**  $h(N) \geq 0$  estimates the cost to go from  $STATE(N)$  to a goal state

Its value is **independent of the current search tree**; it depends only on  $STATE(N)$  and the goal test  $GOAL?$

- Example:

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

Goal state

$h_1(N)$  = number of misplaced numbered tiles = 6

[Why is it an estimate of the distance to the goal?]

# Other Examples

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	


Goal state

- $h_1(N)$  = number of misplaced numbered tiles = 6
- $h_2(N)$  = sum of the (Manhattan) distance of every numbered tile to its goal position  
=  $2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$
- $h_3(N)$  = sum of permutation inversions  
=  $n_5 + n_8 + n_4 + n_2 + n_1 + n_7 + n_3 + n_6$   
=  $4 + 6 + 3 + 1 + 0 + 2 + 0 + 0$   
= 16



# 8-Puzzle

$f(N) = h(N)$  = number of misplaced numbered tiles




The white tile is the empty tile

# 8-Puzzle


$$f(N) = g(N) + h(N)$$

with  $h(N)$  = number of misplaced numbered tiles




# 8-Puzzle

$f(N) = h(N) = \sum$  distances of numbered tiles to their goals



# Robot Navigation




$$h_1(N) = \sqrt{(x_N - x_g)^2 + (y_N - y_g)^2}$$

( $L_2$  or Euclidean distance)

$$h_2(N) = |x_N - x_g| + |y_N - y_g|$$

( $L_1$  or Manhattan distance)

# Best-First $\nrightarrow$ Efficiency



$f(N) = h(N) = \text{straight distance to the goal}$

# Can we prove anything?

- If the state space is infinite, in general the search is not complete
- If the state space is finite and we do not discard nodes that revisit states, in general the search is not complete
- If the state space is finite and we discard nodes that revisit states, the search is complete, but in general is not optimal

# Admissible Heuristic

- Let  $h^*(N)$  be the cost of the optimal path from  $N$  to a goal node
- The heuristic function  $h(N)$  is **admissible** if:

$$0 \leq h(N) \leq h^*(N)$$

- An admissible heuristic function is always **optimistic** !

# Admissible Heuristic

- Let  $h^*(N)$  be the cost of the optimal path from  $N$  to a goal node
- The heuristic function  $h(N)$  is **admissible** if:

$$0 \leq h(N) \leq h^*(N)$$

- An admissible heuristic function is always **optimistic** !

$G$  is a goal node  $\Rightarrow h(G) = 0$



# 8-Puzzle Heuristics

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$  = number of misplaced tiles = 6  
is ???

# 8-Puzzle Heuristics

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$  = number of misplaced tiles = 6  
is **admissible**
- $h_2(N)$  = sum of the (Manhattan) distances of every tile to its goal position  
=  $2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$   
is **???**

# 8-Puzzle Heuristics

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$  = number of misplaced tiles = 6  
is admissible
- $h_2(N)$  = sum of the (Manhattan) distances of every tile to its goal position  
=  $2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$   
is admissible
- $h_3(N)$  = sum of permutation inversions  
=  $4 + 6 + 3 + 1 + 0 + 2 + 0 + 0 = 16$   
is ???

# 8-Puzzle Heuristics

5		8
4	2	1
7	3	6


STATE(N)

1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$  = number of misplaced tiles = 6  
is admissible
- $h_2(N)$  = sum of the (Manhattan) distances of every tile to its goal position  
=  $2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$   
is admissible
- $h_3(N)$  = sum of permutation inversions  
=  $4 + 6 + 3 + 1 + 0 + 2 + 0 + 0 = 16$   
is **not admissible**


# Robot Navigation Heuristics



Cost of one horizontal/vertical step = 1  
Cost of one diagonal step =  $\sqrt{2}$

$$h_1(N) = \sqrt{(x_N - x_g)^2 + (y_N - y_g)^2} \text{ is admissible}$$


# Robot Navigation Heuristics



Cost of one horizontal/vertical step = 1  
Cost of one diagonal step =  $\sqrt{2}$

$$h_2(N) = |x_N - x_g| + |y_N - y_g| \quad \text{is } ???$$

# Robot Navigation Heuristics




Cost of one horizontal/vertical step = 1

Cost of one diagonal step =  $\sqrt{2}$

$h_2(N) = |x_N - x_g| + |y_N - y_g|$  is **admissible** if moving along diagonals is not allowed, and **not admissible** otherwise

$$h^*(I) = 4\sqrt{2}$$

$$h_2(I) = 8$$



# How to create an admissible $h$ ?

- An admissible heuristic can usually be seen as the cost of an optimal solution to a **relaxed** problem (one obtained by removing constraints)
- In robot navigation:
  - The Manhattan distance corresponds to removing the obstacles
  - The Euclidean distance corresponds to removing both the obstacles and the constraint that the robot moves on a grid
- More on this topic later



# A\* Search

(most popular algorithm in AI)

1)  $f(N) = g(N) + h(N)$ , where:

- $g(N)$  = cost of best path found so far to  $N$
- $h(N)$  = **admissible** heuristic function

2) for all arcs:  $c(N, N') \geq \epsilon > 0$

3) **SEARCH#2** algorithm is used

➔ Best-first search is then called **A\* search**


# Result #1

$A^*$  is complete and optimal

[This result holds if nodes revisiting states are not discarded]

# Proof (1/2)


1) If a solution exists, A\* terminates and returns a solution



- For each node  $N$  on the fringe,  
 $f(N) = g(N) + h(N) \geq g(N) \geq d(N) \times \epsilon$ ,  
where  $d(N)$  is the depth of  $N$  in the tree

# Proof (1/2)


## 1) If a solution exists, A\* terminates and returns a solution



- For each node  $N$  on the fringe,  
 $f(N) = g(N) + h(N) \geq g(N) \geq d(N) \times \epsilon$ ,  
where  $d(N)$  is the depth of  $N$  in the tree
- As long as A\* hasn't terminated, a node  $K$  on the fringe lies on a solution path

# Proof (1/2)

## 1) If a solution exists, A\* terminates and returns a solution




- For each node  $N$  on the fringe,  
 $f(N) = g(N) + h(N) \geq g(N) \geq d(N) \times \epsilon$ ,  
where  $d(N)$  is the depth of  $N$  in the tree
- As long as A\* hasn't terminated, a node  $K$  on the fringe lies on a solution path
- Since each node expansion increases the length of one path,  $K$  will eventually be selected for expansion, unless a solution is found along another path

# Proof (2/2)

## 2) Whenever $A^*$ chooses to expand a goal node, the path to this node is optimal

- $C^*$  = cost of the optimal solution path
- $G'$ : non-optimal goal node in the fringe  
 $f(G') = g(G') + h(G') = g(G') > C^*$
- A node  $K$  in the fringe lies on an optimal path:  
$$f(K) = g(K) + h(K) \leq C^*$$
- So,  $G'$  will not be selected for expansion



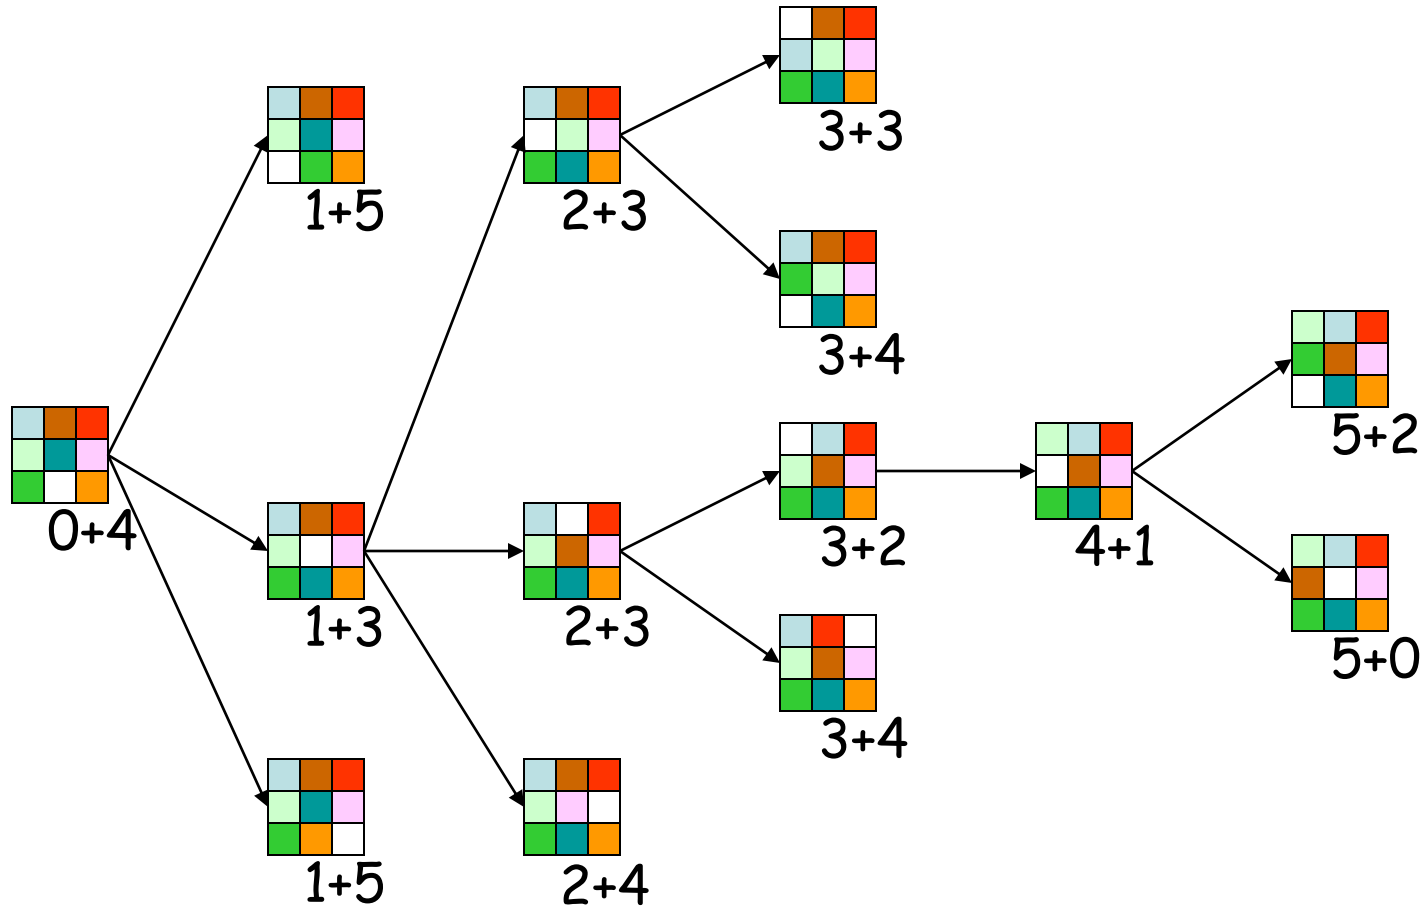
# Time Limit Issue

- When a problem has no solution,  $A^*$  runs for ever if the state space is infinite or states can be revisited an arbitrary number of times. In other cases, it may take a huge amount of time to terminate
- So, in practice,  $A^*$  is given a time limit. If it has not found a solution within this limit, it stops. Then there is no way to know if the problem has no solution, or if more time was needed to find it
- When AI systems are “small” and solving a single search problem at a time, this is not too much of a concern.
- When AI systems become larger, they solve many search problems concurrently, **some with no solution.**  
**What should be the time limit for each of them?**  
More on this in the lecture on Motion Planning ...

# 8-Puzzle


$$f(N) = g(N) + h(N)$$

with  $h(N)$  = number of misplaced tiles





# Robot Navigation



# Robot Navigation

$f(N) = h(N)$ , with  $h(N)$  = Manhattan distance to the goal  
(not  $A^*$ )

8	7	6	5	4	3	2	3	4	5	6
7		5	4	3						5
6			3	2	1	0	1	2		4
7	6									5
8	7	6	5	4	3	2	3	4	5	6

# Robot Navigation

$f(N) = h(N)$ , with  $h(N)$  = Manhattan distance to the goal  
(not  $A^*$ )

8	7	6	5	4	3	2	3	4	5	6
7		5	4	3						5
6			3	2	1	0	1	2		4
7	6									5
8	7	6	5	4	3	2	3	4	5	6

# Robot Navigation

$f(N) = g(N) + h(N)$ , with  $h(N)$  = Manhattan distance to goal  
( $A^*$ )

8+3	7+4	6+3	5+6	4+7	3+8	2+9	3+10	4	5	6
7+2		5+6	4+7	3+8						5
6+1			3	2+9	1+10	0+11	1	2		4
7+0	6+1									5
8+1	7+2	6+3	5+4	4+5	3+6	2+7	3+8	4	5	6

# Best-First Search

- An **evaluation function**  $f$  maps each node  $N$  of the search tree to a real number  $f(N) \geq 0$
- **Best-first search** sorts the FRINGE in increasing  $f$

# A\* Search

1)  $f(N) = g(N) + h(N)$ , where:

- $g(N)$  = cost of best path found so far to  $N$
- $h(N)$  = **admissible** heuristic function

2) for all arcs:  $c(N, N') \geq \epsilon > 0$

3) **SEARCH#2** algorithm is used


➔ Best-first search is then called **A\*** search

# Result #1

$A^*$  is complete and optimal

[This result holds if nodes revisiting states are not discarded]



# What to do with revisited states?



The heuristic  $h$  is clearly admissible





# What to do with revisited states?



If we discard this new node, then the search algorithm expands the goal node next and returns a non-optimal solution



# What to do with revisited states?



Instead, if we do not discard nodes revisiting states, the search terminates with an optimal solution


# But ...

If we do not discard nodes revisiting states, the size of the search tree can be exponential in the number of visited states




# But ...

If we do not discard nodes revisiting states, the size of the search tree can be exponential in the number of visited states



$2n+1$  states




$O(2^n)$  nodes

- It is not harmful to discard a node revisiting a state **if the cost of the new path to this state is  $\geq$  cost of the previous path**  
[so, in particular, one can discard a node if it re-visits a state already visited by one of its ancestors]
- A\* remains optimal, but states can still be re-visited multiple times  
[the size of the search tree can still be exponential in the number of visited states]
- Fortunately, for a large family of admissible heuristics - **consistent** heuristics - there is a much more efficient way to handle revisited states

# Consistent Heuristic

An admissible heuristic  $h$  is **consistent** (or **monotone**) if for each node  $N$  and each child  $N'$  of  $N$ :


$$h(N) \leq c(N, N') + h(N')$$



→ Intuition: a consistent heuristics becomes more precise as we get deeper in the search tree

# Consistency Violation

If  $h$  tells that  $N$  is 100 units from the goal, then moving from  $N$  along an arc costing 10 units should **not** lead to a node  $N'$  that  $h$  estimates to be 10 units away from the goal



# Consistent Heuristic

(alternative definition)

A heuristic  $h$  is **consistent** (or **monotone**) if


1) for each node  $N$  and each child  $N'$  of  $N$ :

$$h(N) \leq c(N, N') + h(N')$$

2) for each goal node  $G$ :

$$h(G) = 0$$

A consistent heuristic is also admissible





# Admissibility and Consistency

- A consistent heuristic is also admissible
- An admissible heuristic may not be consistent, but many admissible heuristics are consistent

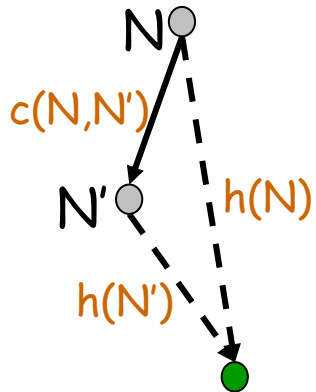
# 8-Puzzle

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

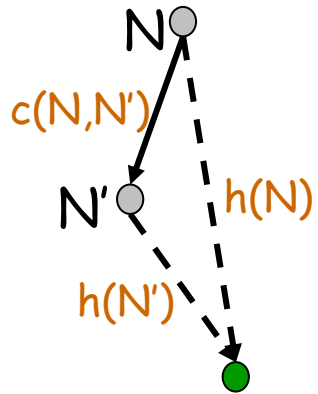
goal



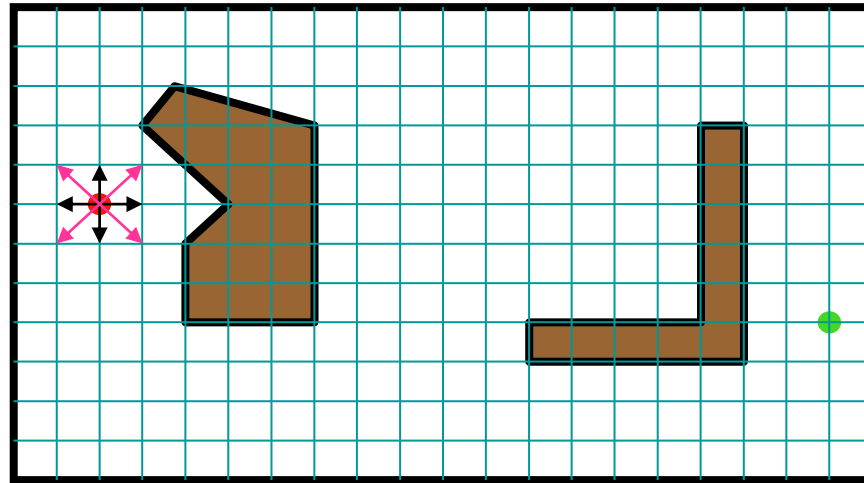
$$h(N) \leq c(N, N') + h(N')$$

- $h_1(N)$  = number of misplaced tiles
  - $h_2(N)$  = sum of the (Manhattan) distances of every tile to its goal position
- are both consistent (why?)

# Robot Navigation



$$h(N) \leq c(N, N') + h(N')$$



Cost of one horizontal/vertical step = 1  
 Cost of one diagonal step =  $\sqrt{2}$

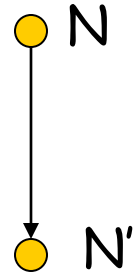
$$h_1(N) = \sqrt{(x_N - x_g)^2 + (y_N - y_g)^2} \quad \text{is consistent}$$

$$h_2(N) = |x_N - x_g| + |y_N - y_g| \quad \text{is consistent if moving along diagonals is not allowed, and not consistent otherwise}$$

## Result #2

If  $h$  is consistent, then whenever  $A^*$  expands a node, it has already found an optimal path to this node's state

# Proof (1/2)



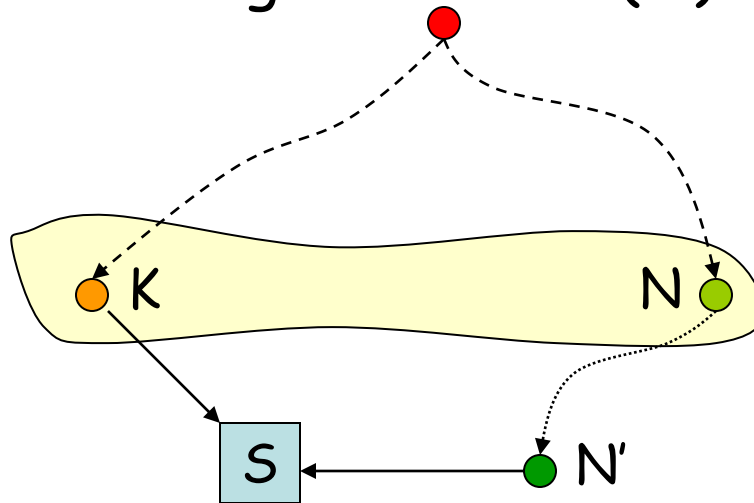
- 1) Consider a node  $N$  and its child  $N'$   
Since  $h$  is consistent:  $h(N) \leq c(N, N') + h(N')$

$$f(N) = g(N) + h(N) \leq g(N) + c(N, N') + h(N') = f(N')$$

So,  $f$  is **non-decreasing** along any path

## Proof (2/2)

- 2) If a node  $K$  is selected for expansion, then any other node  $N$  in the fringe verifies  $f(N) \geq f(K)$



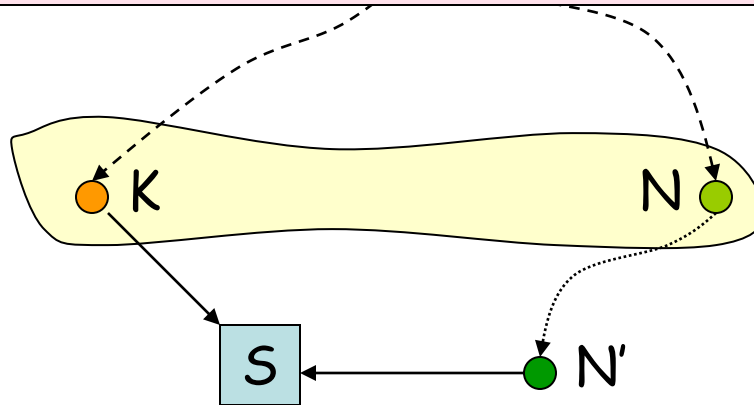
If one node  $N$  lies on another path to the state of  $K$ , the cost of this other path is no smaller than that of the path to  $K$ :

$$f(N') \geq f(N) \geq f(K) \quad \text{and} \quad h(N') = h(K)$$

$$\text{So, } g(N') \geq g(K)$$

## Result #2

If  $h$  is consistent, then whenever  $A^*$  expands a node, it has already found an optimal path to this node's state




If one node  $N$  lies on another path to the state of  $K$ , the cost of this other path is no smaller than that of the path to  $K$ :

$$f(N') \geq f(N) \geq f(K) \quad \text{and} \quad h(N') = h(K)$$

$$\text{So, } g(N') \geq g(K)$$

# Implication of Result #2





# Revisited States with Consistent Heuristic

- When a node is expanded, store its state into CLOSED
- When a new node  $N$  is generated:
  - If  $STATE(N)$  is in CLOSED, discard  $N$
  - If there exists a node  $N'$  in the fringe such that  $STATE(N') = STATE(N)$ , discard the node -  $N$  or  $N'$  - with the largest  $f$  (or, equivalently,  $g$ )

Is  $A^*$  with some consistent heuristic all that we need?

No !

There are **very dumb** consistent heuristic functions

For example:  $h \equiv 0$

- It is consistent (hence, admissible) !
- $A^*$  with  $h \equiv 0$  is uniform-cost search
- Breadth-first and uniform-cost are particular cases of  $A^*$

# Heuristic Accuracy

Let  $h_1$  and  $h_2$  be two consistent heuristics such that for all nodes  $N$ :

$$h_1(N) \leq h_2(N)$$

$h_2$  is said to be **more accurate** (or **more informed**) than  $h_1$

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$  = number of misplaced tiles
- $h_2(N)$  = sum of distances of every tile to its goal position
- $h_2$  is more accurate than  $h_1$

## Result #3

- Let  $h_2$  be more accurate than  $h_1$
- Let  $A_1^*$  be  $A^*$  using  $h_1$   
and  $A_2^*$  be  $A^*$  using  $h_2$
- Whenever a solution exists, all the nodes expanded by  $A_2^*$ , except possibly for some nodes such that
$$f_1(N) = f_2(N) = C^* \text{ (cost of optimal solution)}$$
are also expanded by  $A_1^*$

# Proof

- $C^* = h^*(\text{initial-node})$  [cost of optimal solution]
- Every node  $N$  such that  $f(N) < C^*$  is eventually expanded. No node  $N$  such that  $f(N) > C^*$  is ever expanded
- Every node  $N$  such that  $h(N) < C^* - g(N)$  is eventually expanded. So, every node  $N$  such that  $h_2(N) < C^* - g(N)$  is expanded by  $A_2^*$ . Since  $h_1(N) \leq h_2(N)$ ,  $N$  is also expanded by  $A_1^*$
- If there are several nodes  $N$  such that  $f_1(N) = f_2(N) = C^*$  (such nodes include the optimal goal nodes, if there exists a solution),  $A_1^*$  and  $A_2^*$  may or may not expand them in the same order (until one goal node is expanded)

# Effective Branching Factor

- It is used as a measure the effectiveness of a heuristic
- Let  $n$  be the total number of nodes expanded by  $A^*$  for a particular problem and  $d$  the depth of the solution
- The **effective branching** factor  $b^*$  is defined by  $n = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$

# Experimental Results

(see R&N for details)


- 8-puzzle with:
  - $h_1$  = number of misplaced tiles
  - $h_2$  = sum of distances of tiles to their goal positions
- Random generation of many problem instances
- Average effective branching factors (number of expanded nodes):

d	IDS	$A_1^*$	$A_2^*$
2	2.45	1.79	1.79
6	2.73	1.34	1.30
12	2.78 (3,644,035)	1.42 (227)	1.24 (73)
16	--	1.45	1.25
20	--	1.47	1.27
24	--	1.48 (39,135)	1.26 (1,641)



# How to create good heuristics?

- By solving **relaxed** problems at each node
- In the 8-puzzle, the sum of the distances of each tile to its goal position ( $h_2$ ) corresponds to solving 8 simple problems:



$d_i$  is the length of the shortest path to move tile  $i$  to its goal position, ignoring the other tiles, e.g.,  $d_5 = 2$


$$h_2 = \sum_{i=1, \dots, 8} d_i$$

- It ignores negative interactions among tiles





# Can we do better?

- For example, we could consider two more complex relaxed problems:

$d_{1234}$  = length of the shortest path to move tiles 1, 2, 3, and 4 to their goal positions, ignoring the other tiles



$d_{5678}$




- $h = d_{1234} + d_{5678}$  [disjoint pattern heuristic]





# Can we do better?

- For example, we could consider two more complex relaxed problems:

$d_{1234}$  = length of the shortest path to move tiles 1, 2, 3, and 4 to their goal positions, ignoring the other tiles



$d_{5678}$




- $h = d_{1234} + d_{5678}$  [disjoint pattern heuristic]
- How to compute  $d_{1234}$  and  $d_{5678}$ ?

# Can we do better?

- For example, we could consider two more complex relaxed problems:


$d_{1234}$  = length of the shortest path to move tiles 1, 2, 3, and 4 to their goal positions, ignoring the other tiles



- $h = d_{1234} + d_{5678}$  [disjoint pattern heuristic]
- These distances are pre-computed and stored  
[Each requires generating a tree of 3,024 nodes/states (breadth-68 first search)]

# Can we do better?

- For example, we could consider two more complex relaxed problems:



- $h = d_{1234} + d_{5678}$  [disjoint pattern heuristic]
- These distances are pre-computed and stored [Each requires generating a tree of 3,024 nodes/states (breadth-first search)]

# On Completeness and Optimality

- $A^*$  with a consistent heuristic function has nice properties: completeness, optimality, no need to revisit states
- Theoretical completeness does not mean “practical” completeness if you must wait too long to get a solution (remember the time limit issue)
- So, if one can't design an accurate consistent heuristic, it may be better to settle for a non-admissible heuristic that “works well in practice”, even though completeness and optimality are no longer guaranteed



# Iterative Deepening A\* (IDA\*)

- Idea: Reduce memory requirement of A\* by applying cutoff on values of  $f$
- Consistent heuristic function  $h$
- Algorithm IDA\*:
  1. Initialize cutoff to  $f(\text{initial-node})$
  2. Repeat:
    - a. Perform depth-first search by expanding all nodes  $N$  such that  $f(N) \leq \text{cutoff}$
    - b. Reset cutoff to smallest value  $f$  of non-expanded (leaf) nodes

# 8-Puzzle

$$f(N) = g(N) + h(N)$$

with  $h(N)$  = number of misplaced tiles






# 8-Puzzle

$$f(N) = g(N) + h(N)$$


with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$


with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$


with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$

with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$

with  $h(N)$  = number of misplaced tiles




Cutoff=5



# 8-Puzzle

$$f(N) = g(N) + h(N)$$



with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$


with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$

with  $h(N)$  = number of misplaced tiles






# 8-Puzzle

$$f(N) = g(N) + h(N)$$


with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$


with  $h(N)$  = number of misplaced tiles



# 8-Puzzle

$$f(N) = g(N) + h(N)$$

with  $h(N)$  = number of misplaced tiles



# Advantages/Drawbacks of IDA\*

- Advantages:
  - Still complete and optimal
  - Requires less memory than  $A^*$
  - Avoid the overhead to sort the fringe
- Drawbacks:
  - Can't avoid revisiting states not on the current path
  - Available memory is poorly used  
(→ memory-bounded search, see R&N p. 101-104)

# Local Search

- Light-memory search method
- No search tree; only the current state is represented!
- Only applicable to problems where the path is irrelevant (e.g., 8-queen), unless the path is encoded in the state
- Many similarities with optimization techniques

# Steepest Descent




- 1)  $S \leftarrow$  initial state
- 2) Repeat:
  - a)  $S' \leftarrow \arg \min_{S' \in \text{SUCCESSORS}(S)} \{h(S')\}$
  - b) if  $\text{GOAL?}(S')$  return  $S'$
  - c) if  $h(S') < h(S)$  then  $S \leftarrow S'$  else return failure

Similar to:

- hill climbing with  $-h$
- gradient descent over continuous space

# Application: 8-Queen

- 1) Pick an initial state  $S$  at random with one queen in each column
- 2) Repeat  $k$  times:
  - a) If  $GOAL?(S)$  then return  $S$
  - b) Pick an attacked queen  $Q$  at random
  - c) Move  $Q$  in its column to minimize the number of attacking queens  $\rightarrow$  new  $S$  [min-conflicts heuristic]
- 3) Return failure



# Application: 8-Queen

Why does it work ???

- 1) There are **many** goal states that are well-distributed over the state space
- 2) If no solution has been found after a few steps, it's better to start it all over again. Building a search tree would be much less efficient because of the high branching factor
- 3) Running time almost independent of the number of queens

				2			
				2			
				3			

							2
							2
							2




# Steepest Descent

- 1)  $S \leftarrow$  initial state
- 2) Repeat:
  - a)  $S' \leftarrow \arg \min_{S' \in \text{SUCCESSORS}(S)} \{h(S')\}$
  - b) if  $\text{GOAL?}(S')$  return  $S'$
  - c) if  $h(S') < h(S)$  then  $S \leftarrow S'$  else return failure

may easily get stuck in local minima

- Random restart (as in n-queen example)
- Monte Carlo descent

# Monte Carlo Descent

- 1)  $S \leftarrow$  initial state
- 2) Repeat  $k$  times:
  - a) If  $GOAL?(S)$  then return  $S$
  - b)  $S' \leftarrow$  successor of  $S$  picked at random
  - c) if  $h(S') \leq h(S)$  then  $S \leftarrow S'$
  - d) else
    - $\Delta h = h(S') - h(S)$
    - with probability  $\sim \exp(-\Delta h/T)$ , where  $T$  is called the "temperature", do:  $S \leftarrow S'$  [Metropolis criterion]
- 3) Return failure

**Simulated annealing** lowers  $T$  over the  $k$  iterations.  
It starts with a large  $T$  and slowly decreases  $T$


# “Parallel” Local Search Techniques

They perform several local searches concurrently, but not independently:


- Beam search
- Genetic algorithms

See R&N, pages 115-119


Search problems



Blind search




Heuristic search:  
best-first and  $A^*$



Construction of heuristics



Variants of  $A^*$



Local search

# When to Use Search Techniques?

- 1) The search space is small, and
  - No other technique is available, or
  - Developing a more efficient technique is not worth the effort
- 2) The search space is large, and
  - No other available technique is available, and
  - There exist “good” heuristics